



Hive

THE PORTAL

THE APIARY · LEGION CODE × ACTIVELOOP

How do your engineers see the health of their AI stack without port-hunting?

The Apiary runs several local services per developer, and a status view scattered across loopback ports is a status view nobody checks. Hive is one always-on portal at 127.0.0.1:3853 that serves the entire dashboard and a live health rail, built to be the last thing standing when something else goes down.

Why does a single portal matter operationally?

Because visibility that takes effort gets skipped. Hive puts memories, graphs, sync, logs, ROI, and per-service health behind one address, so an engineer always knows the state of their stack at a glance.

Does it become a new point of failure?

No. Hive is its own service, booted at startup and watched by Doctor, and if a service behind it drops, that panel says so while the rest keeps working and recovers on its own. It was built to survive the exact moment you need a status view.

What is its security surface?

Minimal. It binds to loopback only, so nothing off-device can reach it, and it stores nothing: your session passes straight through to the services that own your data. Sign-in creates your Deeplake credential on your machine, not in the portal.

What does it give an operator that per-service pages don't?

One honest picture. A health rail on every page, a readiness screen on cold boot, and per-service metrics and logs in one place, so triage starts from a single source of truth instead of a tab hunt.