



Honeycomb

THE MEMORY

THE APIARY · LEGION CODE × ACTIVELOOP

Why do you re-explain your own project to your AI every morning?

You taught your assistant your stack and the fix that finally worked at midnight, then closed the window and it forgot all of it. Honeycomb is the memory it never had: install once and it hands the right notes back to any assistant, on any machine. Learn something once, recall it everywhere.

I told Cursor yesterday. Why does Claude Code have no idea today?

Their memory dies with the window, and they do not talk to each other. Honeycomb sits underneath all of them (Claude Code, Cursor, and Codex today) and gives them one shared brain, so a note written in one is recalled by another, even on a different laptop.

Do I have to babysit it?

No. Install opens the Hive dashboard, where you connect your assistants in a couple of clicks, and from then on it captures the useful moments on its own and briefs every new session. Your agent starts smart instead of blank.

How does it find the right memory when I forget the exact words?

Your memory lives on Deeplake, which searches by meaning and not just keywords, so it surfaces the note you needed even when you would never have guessed the term. It also tidies itself over time, so the more you use it the sharper it gets, not the noisier.

Is my code safe?

Yes. Your memory sits on Deeplake, reached only by the local helper on your machine and never by the assistant directly, with secrets never shown to the agent. Working on something sensitive? Flip to read-only and nothing new is written, while recall still works.