



Honeycomb

THE MEMORY

THE APIARY · LEGION CODE × ACTIVELOOP

Why pay to teach the same context to AI every single day?

Your developers' assistants forget everything the moment a session ends, so the team re-explains the same codebase and re-buys the same tokens every morning. Honeycomb gives every agent one shared, lasting memory on Deeplake: what one learns, all recall, and it compounds instead of evaporating.

What does a cold start actually cost us?

Every fresh session re-sends context you pay for and re-explains what the tool knew yesterday. Honeycomb primes each session up front, and the ROI page shows real-dollar savings, marked measured or estimated, so it reads like a receipt and not a guess.

Is our memory secure?

It lives on Deeplake, reached only by a small local daemon and never by the assistant itself. Teams and projects are isolated at the storage layer, secrets are never shown to an agent, and every memory is versioned, so you always know what was known and when.

Why Deeplake, and not a bolt-on vector store?

Deeplake holds your memory as exact text and as meaning at once, so recall finds the right note even when nobody used the right words, at scale, with full history behind every fact. It grows more powerful and more auditable, instead of rotting into a cache.

How does one engineer's fix reach the whole team?

A solved problem becomes a skill that appears automatically in every teammate's assistant on their next session, with no wiki and no copy-paste. Sharing is opt-in, so private stays private until someone widens it on purpose.